

Run #133 — confirmation desk day: the first vision board

Dates: started 2026-08-29, completed 2026-08-30 — the leaderboard dates a board by when its run was created, so it files this one under the 29th **Workflow:** confirmation-desk-day , 9 steps / 33 checks, **uncalibrated par Contestants:** the four vision-tagged models, **2 trials each**, effort and output budget unpinned (32768 tokens, both wire protocols) **Run:** #133 on the live desk database · 8 clean trials · 0 invalid · 0 errors

Board

#	Contestant	Objective	OVR	GRD	ADH	SYN	PRC	EFF	CON	calls
1	deepseek-v4-flash-vision	86.3	81	94	99	99	60	24	96	46,30
2	glm-5-3-flash	87.9	78	89	99	99	70	18	86	63,37
3	gemini-3-7-flash	81.8	75	84	99	99	54	16	89	47,59
4	gpt-5-6-luna	84.8	68	59	99	86	94	10	89	59,74

Zero truncation, zero malformed tool calls, zero step errors across all eight trials.

The highest objective score does not win the board. glm-5-3-flash posts 87.9 to deepseek's 86.3 and still places second: grounding carries 0.32 of OVR and deepseek reads better (94 vs 89), consistently (CON 96 vs 86), in fewer calls. gpt-5-6-luna shows the same effect at the bottom — 84.8 objective, third-highest in the field, but last on OVR because its GRD is 59.

What this board was built to answer

Can GLM 5.3 Flash — z.ai's first multimodal model — work a confirmations desk against Gemini, GPT and DeepSeek? Yes, on correctness: it records the highest objective score in the

field and ties the leaders on every designed vision trap it was given. It places second because it is neither the most consistent nor the leanest.

The more useful answer is that **the question is nearly obsolete at this tier.**

The designed vision traps are solved

Three of the four traps this workflow was built around are **saturated at 8/8** — every contestant, every trial:

Trap	Document	Truth	Rate
Image-only scan, no text layer	conf-04 GOOGL	strike 205.0	8/8
Struck-through and amended strike	conf-09 NVDA	1045.0, not the struck value	8/8
Barrier direction carried only by a ticked box	conf-10 AMZN	DOWN_OUT , 171.2	8/8
A term the document never states	conf-07 META	initial_price null, status invalid	8/8

That last row is the one worth pausing on. Step 7 asks for a value that is not in the document, with the strike sitting nearby as an obvious substitute. **No model substituted, in any trial.** The pilot board saw the incumbent substitute in 2 of 6 runs; on this field the behaviour is gone.

What still discriminates — 6 of 28 checks

Per-check pass-rate tally across the whole field. Keyed by label, so a check appearing in several steps is counted once per occurrence.

Check	Rate	What it is
Step 6 — ORCL notional 636,000	2/8	A low-contrast number beside a decoy. The hardest thing on the board.
<code>get_confirmation_batch</code> called	24/48	Procedure, not sight
Step 8 — <code>skipped_count</code> = 1	4/8	Booking restraint and counting
Step 3 — AMD strike 185.0	6/8	Priced terms on page 2 of a mixed text-and-scan document
Step 3 — AMD initial_price 178.9	6/8	Same document, same page
Step 9 — artifact names ARD-EQO-2026-04901	7/8	Carrying a reference into the summary

22 of 28 checks are dead (0/N or N/N). This tally is only valid because every arm on this board is healthy — a contaminated arm fails checks the healthy field passes and manufactures apparent discrimination. On the pilot board, with one arm poisoned by an infra defect, the same instrument read 8 checks dead where the repaired board read 25.

The two things that are genuinely hard are not OCR

A faint number is harder than a degraded page. The ORCL notional is rendered at low contrast next to a decoy value. It is not a scan, not handwriting, not a checkbox — it is ordinary printed text that is merely *faint*, and it is the only check on the board that most of the field misses. Two of eight trials read it: one gemini trial and one deepseek trial. **No model read it in both of its trials.** GLM and luna missed it 0/2 each.

A second page is harder than a second modality. conf-08 is the mixed text-and-scan document, and its priced terms are on page 2. The step's prompt says so explicitly. Both of its checks come in at 6/8 — and both misses are the same contestant, `gpt-5-6-luna`, in both trials.

The board's variance is procedural, not visual

`glm-5-3-flash` scored **97.0 and then 78.8** — an 18-point spread on identical inputs, and the reason it dropped from first to second when its second trial landed.

Every point of that spread is one check. Its two trials read the documents *identically*: same strikes, same barrier, same absent term, same faint-notional miss. What differed is that trial 1 called `get_confirmation_batch` at all six graded steps and trial 2 called it at none.

`deepseek-v4-flash-vision` shows the same thing inverted. Its lower-scoring trial (81.8, 30 calls) is the one where it read **every trap correctly including the faint notional** — the only trial on the board to do so — and it scores lower than its 90.9 trial because it skipped the batch re-reads.

So the single widest-swinging check on this board rewards re-fetching a payload the model already holds. That is the concern the pilot report raised about these checks, now with a number on it: `get_confirmation_batch` **accounts for more objective variance than every vision check combined**. It measures procedure, and a model answering correctly from context it already has is arguably being efficient rather than wrong. This is the same class as the "step-scoped check penalises reading the evidence one step early" defect already recorded in `CLAUDE.md`.

Par is still not calibrated, and now there is evidence for why

`confirmation-desk-day` declares no `par_tool_calls`, so EFF runs on the legacy hyperbolic curve against a `designed_par` of 9 — the theoretical minimum. That crushes every EFF on this board into 10–24 on a 0–99 scale, and EFF is 0.16 of OVR.

The established way to calibrate is the median of the **fully-correct** trials. **This board produced none: not one of the eight trials passed every check.** The pilot's single 33/33 run is one data point, not a distribution. So par stays unset rather than being invented from a sample that does not exist, and **EFF and OVR here are not comparable with calibrated boards.**

For the record, the counted call distribution over 8 clean trials: median **53**, range **30–74**.

Why this board can be trusted

The pilot board found three harness defects, each of which cost a real contestant real points, and each invisible to the golden replay. All three are fixed here, and this run is the evidence:

1. **StreamChunkTimeoutError after 120s.** Step 1 parses six documents in one tool body — 12+ vision calls synchronously — so the outer stream emits nothing for minutes and langchain's default fires on a connection that is idle rather than dead. Launched with `LANGCHAIN_OPENAI_STREAM_CHUNK_TIMEOUT_S=900`; zero contestants died to it.

2. **Unaddressable upload paths.** A model told the files are at `/artifacts/uploads/confirmations/` and then rejected for using that spelling scored 30.3 with grounding 0/10 on the pilot. Fixed; zero path errors here.
3. **A binary `read_file` poisons the conversation permanently.** `deepagents` returns a binary read as a media content block; three of four gateways reject that shape from a *tool* message, each in its own dialect, and the rejected message stays in the history so every later turn redraws the same 400. On the pilot this took `deepseek-v4-flash-vision` from a real 100.0 to a recorded 36.4 and last place. `BinaryReadGuardMiddleware` now intercepts it in all three hand-built agent stacks **and** in the `general-purpose` subagent that `deepagents` adds for itself — a fourth stack that inherited the parent's full toolset while running unguarded and unaudited.

Zero errors across eight trials is the result of those three fixes, not of an easy workflow.

Caveats a reader must carry

- **Do not read this ordering as a vision ranking.** Three of four contestants score GRD 84–94 and 22 of 28 checks are saturated. Separation comes from consistency, procedure and volume.
- **EFF is uncalibrated** (see above). Do not compare it across boards.
- **One arm was re-run.** `glm-5-3-flash` lost a trial to a transport-level `RemoteProtocolError` — the peer closed the connection mid-body — and infra trials are skipped rather than retried, which would have left the headline contestant at half the field's depth with no CON. Its `scored` row was deleted and the arm re-run to two fresh trials via `--resume 133`, at the same effort and the same re-supplied output budget. **The repair changed the podium:** on the single surviving trial GLM stood at OVR 85 and first place; at full depth it is 78 and second. Its objective mean is 87.9 either way — the coincidence is real and is why OVR, not the objective percentage, ranks the board.
- **Two trials, not five.** CON is measured but from a small sample.
- **The extraction sub-call routes to the contestant.** Without that, every model would read every document with one shared extractor and every vision check would land N/N carrying no signal.

Next

- **Reconsider the `get_confirmation_batch` checks.** They are 1 of the 6 surviving discriminators and the largest single source of objective variance, and what they reward is re-fetching held context.

- **The workflow needs harder traps.** Its designed vision difficulty is solved by the whole field. The two things that still separate contestants — low contrast and a second page — were incidental properties of the corpus, not designed ones. That is where the next iteration should aim.
- **Do not calibrate par until a trial passes every check.** The instrument is the median of fully-correct trials, and that set is currently empty.