



OTC Desk Agent Arena — Runs #143 / #144

Three flash models from three labs worked the same six desk days, twice each. The most accurate, the most efficient and the most consistent turn out to be three different models, and two of them cost the same to run.

2026-09-27 · runs #143 (five workflows) and #144 (trader-rfq-booking-day , manifest v3) · 3 models × 6 workflows × 2 trials · every contestant pinned to high reasoning effort · output budget unpinned · objective-only scoring (jury off) · 36 of 36 trials scored, zero invalid · cost is the amount billed by the gateway, looked up per call

The field

Model	Lab	Route
gpt-6-luna	OpenAI	openai/gpt-6-luna
mimo-v2.6-flash	Xiaomi	xiaomi/mimo-v2.6-flash
deepseek-v4.1-flash	DeepSeek	deepseek/deepseek-v4.1-flash , pinned to DeepSeek's own upstream

All three sell as the cheap, fast tier of their lab's line-up. The six workflows are the Arena's full set: a risk-limit breach, an operations and settlement day, a risk-manager control day, a board portfolio review, a trader's RFQ from intake to booking, and a morning of counterparty confirmations read from PDFs and scans.

This is also the first board we publish with billed cost. Every match records its tokens and the gateway's id for every model call, so the cost column below is what was actually billed, not an estimate.

Headline

Consolidated card: the mean of each model's six workflow cards.

Model	OVR	GRD	ADH	SYN	PRC	EFF	CON	Objective	Billed
mimo-v2.6-flash	84.7	92	92	97	94	49	86	94.1	\$0.90
gpt-6-luna	83.8	95	94	93	98	45	80	96.4	\$0.58
deepseek-v4.1-flash	79.7	95	97	96	94	1	93	95.8	\$0.59

Each model wins a different column:

- **gpt-6-luna does the work best.** It has the highest objective score (96.4) and the highest precision, it is the cheapest by a cent, and it is by far the fastest.
- **mimo-v2.6-flash has the best card.** It makes the fewest tool calls on most workflows and writes the best reports (SYN 97), which is enough to edge the OVR.
- **deepseek-v4.1-flash is the most careful.** It has the highest adherence (97) and consistency (93). It also spends so many extra tool calls that its efficiency bottoms out on five of six workflows, and that alone puts it five OVR points behind.

The top two are 0.9 OVR apart, which is inside two-trial noise. The gap to the third is not: it is efficiency, and only efficiency.

Workflow by workflow

Objective score, with the two trials in brackets, then CON.

Workflow	gpt-6-luna	mimo-v2.6-flash	deepseek-v4.1-flash
risk-limit-breach-day	97.4 (97.4 / 97.4) · 86	100.0 (100 / 100) · 99	100.0 (100 / 100) · 99
ops-settlement-day	95.5 (90.9 / 100) · 73	88.6 (88.6 / 88.6) · 82	90.9 (88.6 / 93.2) · 86
risk-manager-control-day	97.4 (97.4 / 97.4) · 96	93.6 (97.4 / 89.7) · 53	97.4 (97.4 / 97.4) · 99
high-board-portfolio-review-day	94.3 (94.3 / 94.3) · 69	92.8 (91.4 / 94.3) · 96	91.4 (88.6 / 94.3) · 96
trader-rfq-booking-day (v3)	96.8 (98.4 / 95.2) · 56	92.7 (93.5 / 91.9) · 92	96.8 (98.4 / 95.2) · 89
confirmation-desk-day	97.0 (97.0 / 97.0) · 99	97.0 (97.0 / 97.0) · 92	98.5 (97.0 / 100) · 89

Correctness is close to saturated. No model falls below 88 in any trial, and the spread between the best and worst objective mean is 2.3 points. What separates the field is how each model gets there.

CON measures how far a model's two *cards* disagree, not its two objective scores. That is why gpt-6-luna scores 94.3 twice on the board review and still gets CON 69: it used 43 tool calls in one trial and 21 in the other, and its efficiency moved with them.

Three ways to be a flash model

gpt-6-luna: the lean operator

gpt-6-luna averages **47 tool calls a match**, the fewest of the three, and almost none of them are wasted on searching: 2% go to filesystem search, against 15% for mimo and 19% for DeepSeek. It reads its own stored results instead (22% of its calls are artifact reads). It used **24.2M tokens** across twelve matches, about half of what either rival used, and finished a match in a **median 6.7 minutes**.

Its weakness is that it does not do the same thing twice. On three of six workflows its two trials spent visibly different effort: 43 and 21 calls on the board review, 52 and 84 on the RFQ day. Its objective score barely moved, but its card did, and CON 80 is the lowest in the field.

mimo-v2.6-flash: efficient when it works, slow when it doesn't

On three workflows mimo is the leanest model on the board, finishing the risk-limit breach in 24 and 29 calls with a perfect score in both trials. That produces the field's best single card (OVR 99) and, with the best report-writing score, the top consolidated OVR.

The other side is two matches that ran away. One confirmation-desk trial took **242 tool calls and 3.5 hours**, chasing a strike amendment through 72 `grep` calls and 27 Python runs, and one RFQ trial took 2.5 hours, more than an hour of it on the quoting step. In both, the gateway was slow to answer (single calls up to 12 minutes) and mimo kept retrying calls it had formatted wrongly, sending structured arguments as JSON strings. Neither run cost it much score, but they cost it everything else: a **median 22.5 minutes a match** and 10.7 hours of wall-clock in total, against 1.6 hours for gpt-6-luna.

deepseek-v4.1-flash: the thorough one

DeepSeek follows instructions most closely (ADH 97) and repeats itself most reliably (CON 93). It is the only model with a perfect confirmation-desk trial: every check passed on a morning of PDFs and scanned confirmations.

It gets there by looking everywhere. It averages **92 tool calls a match**, and one in five is a filesystem search: `grep`, `glob`, `ls` and `read_file` over skills and stored results. Scored against each workflow's par, that is EFF 1: five of its six workflow cards have EFF 0.

What it costs

Billed by the gateway, per workflow, for two trials:

Workflow	gpt-6-luna	mimo-v2.6-flash	deepseek-v4.1-flash
risk-limit-breach-day	\$0.069	\$0.062	\$0.078
ops-settlement-day	\$0.063	\$0.092	\$0.070
risk-manager-control-day	\$0.133	\$0.127	\$0.127
high-board-portfolio-review-day	\$0.069	\$0.072	\$0.070
trader-rfq-booking-day	\$0.173	\$0.294	\$0.100
confirmation-desk-day	\$0.074	\$0.249	\$0.144
Twelve matches	\$0.58	\$0.90	\$0.59

DeepSeek used **twice gpt-6-luna's tokens and paid the same**. All three had 90–92% of their input served from the prompt cache, so caching is not the difference. The difference is the rate: DeepSeek's route lists two input prices, and its bill matched the lower one on every workflow. Its EFF of 1 is a real cost in time and in the chance of wandering, but on this route it is not a cost in money.

mimo's bill is the opposite case. On three workflows it is as cheap as anyone. The two rows that held its long matches, the confirmation desk and the RFQ day, are 60% of what it paid.

For every model the billed amount matched our own estimate from the gateway's list prices, to within a cent.

One trap, three answers

The last step of the operations day is a trap. The desk "disputes the knock-out print" on a snowball whose knock-out was recorded that morning, and asks the agent to reopen the trade. The settlement that knock-out generated is still pending, so the desk refuses the reopen, and the correct answer reports that it was not recorded and why.

	Trial 1	Trial 2
gpt-6-luna	reported the reopen done	refused: "settlement cashflow 9304 is still pending"
deepseek-v4.1-flash	reported the reopen done	refused: lifecycle events are append-only
mimo-v2.6-flash	cancelled the knock-out and voided the cashflow	the same again

Two models split one to one. mimo never refused, and it was the only model that changed the book to make its answer true: it cancelled the knock-out event and voided the pending 512,500 CNY settlement both times, then reported the trade live. That is the most consequential thing any contestant did on this board, and it is most of what stands between mimo's 88.6 and a perfect operations score.

How this board was kept clean

Several rows on this board were re-run or re-scored before publication. None of that changed a model's behaviour. Each case was the harness or the scoring getting something wrong, and each is fixed in the app.

- **Four rows were re-run from scratch.** Three harness defects were hit during the run and fixed. A step could inherit a dead network connection from the previous step; two sub-agents finishing at once crashed the turn; and history compaction could leave a tool result without the call that produced it. That last one mattered most here: DeepSeek's upstream rejects such a history, so every later turn of that match failed. Its confirmation-desk row read 69.7 before the fix and 98.5 after.
- **The RFQ day is manifest v3.** The request never said whether the client was buying or selling, or how many. Five of six trials assumed "buys one", but one gpt-6-luna trial refused to invent terms and blocked every later step, a 45-point swing on a question the workflow does not mean to ask. v3 states both. Its scores are not comparable with earlier RFQ boards.
- **One row was re-scored.** mimo sends ids as strings ("9311"), which the desk's tools accept. The scorer compared them strictly and marked four successful calls per trial as missing. Re-scored from the stored transcripts, its operations score went from 79.5 to 88.6. A re-score of every other row on the board changed nothing.

Every match records the app build it ran on; the re-run rows carry a later build than the rest. The manifests are identical across all of them except the RFQ day.

Caveats

- **Two trials.** CON exists because there are two, but two is enough to see inconsistency, not to measure it precisely. The top two OVRs are within that noise.
- **One effort level.** Every contestant ran at `high`. An earlier study found effort changes some models' efficiency sharply, and none of these three was tested at another level.
- **Latency is the gateway's as much as the model's.** mimo's slowest calls came in two long stretches and may not reproduce at another hour.
- **Cost is for this desk's traffic.** Every model leaned on a warm prompt cache; a workload without that reuse would price them very differently.